

Homogeneous representation of 2D lines and points

- ▶ A 2D line is defined by $ax + by + c = 0$ i.e. a parametrization $\mathbf{l} = (a, b, c)$.

Homogeneous representation of 2D lines and points

- ▶ A 2D line is defined by $ax + by + c = 0$ i.e. a parametrization $\mathbf{l} = (a, b, c)$.
- ▶ However, $kax + kby + kc = 0$ corresponds to the same line, thus $\mathbf{l} = (ka, kb, kc), \forall k \in \mathbb{R} \setminus \{0\}$

Homogeneous representation of 2D lines and points

- ▶ A 2D line is defined by $ax + by + c = 0$ i.e. a parametrization $\mathbf{l} = (a, b, c)$.
- ▶ However, $kax + kby + kc = 0$ corresponds to the same line, thus $\mathbf{l} = (ka, kb, kc), \forall k \in \mathbb{R} \setminus \{0\}$
- ▶ A 2D point (x, y) lies on a line (a, b, c) if $ax + by + c = 0$.

Homogeneous representation of 2D lines and points

- ▶ A 2D line is defined by $ax + by + c = 0$ i.e. a parametrization $\mathbf{l} = (a, b, c)$.
- ▶ However, $kax + kby + kc = 0$ corresponds to the same line, thus $\mathbf{l} = (ka, kb, kc), \forall k \in \mathbb{R} \setminus \{0\}$
- ▶ A 2D point (x, y) lies on a line (a, b, c) if $ax + by + c = 0$.
- ▶ This may be expressed as $(x, y, 1)^T \cdot (a, b, c) = (x, y, 1)^T \cdot \mathbf{l} = 0$.

Homogeneous representation of 2D lines and points

- ▶ A 2D line is defined by $ax + by + c = 0$ i.e. a parametrization $\mathbf{l} = (a, b, c)$.
- ▶ However, $kax + kby + kc = 0$ corresponds to the same line, thus $\mathbf{l} = (ka, kb, kc), \forall k \in \mathbb{R} \setminus \{0\}$
- ▶ A 2D point (x, y) lies on a line (a, b, c) if $ax + by + c = 0$.
- ▶ This may be expressed as $(x, y, 1)^T \cdot (a, b, c) = (x, y, 1)^T \cdot \mathbf{l} = 0$.
- ▶ $\forall k \in \mathbb{R} \setminus \{0\}, (kx, ky, k)^T \cdot \mathbf{l} = 0$ if and only if $(x, y, 1)^T \cdot \mathbf{l} = 0$.

Homogeneous representation of 2D lines and points

- ▶ A 2D line is defined by $ax + by + c = 0$ i.e. a parametrization $\mathbf{l} = (a, b, c)$.
- ▶ However, $kax + kby + kc = 0$ corresponds to the same line, thus $\mathbf{l} = (ka, kb, kc), \forall k \in \mathbb{R} \setminus \{0\}$
- ▶ A 2D point (x, y) lies on a line (a, b, c) if $ax + by + c = 0$.
- ▶ This may be expressed as $(x, y, 1)^T \cdot (a, b, c) = (x, y, 1)^T \cdot \mathbf{l} = 0$.
- ▶ $\forall k \in \mathbb{R} \setminus \{0\}, (kx, ky, k)^T \cdot \mathbf{l} = 0$ if and only if $(x, y, 1)^T \cdot \mathbf{l} = 0$.
- ▶ $\forall k \in \mathbb{R} \setminus \{0\}$, we denote thus (kx, ky, k) as the homogeneous representation of the 2D point (x, y) .

Homogeneous representation of 2D lines and points

- ▶ A 2D line is defined by $ax + by + c = 0$ i.e. a parametrization $\mathbf{l} = (a, b, c)$.
- ▶ However, $kax + kby + kc = 0$ corresponds to the same line, thus $\mathbf{l} = (ka, kb, kc), \forall k \in \mathbb{R} \setminus \{0\}$
- ▶ A 2D point (x, y) lies on a line (a, b, c) if $ax + by + c = 0$.
- ▶ This may be expressed as $(x, y, 1)^T \cdot (a, b, c) = (x, y, 1)^T \cdot \mathbf{l} = 0$.
- ▶ $\forall k \in \mathbb{R} \setminus \{0\}, (kx, ky, k)^T \cdot \mathbf{l} = 0$ if and only if $(x, y, 1)^T \cdot \mathbf{l} = 0$.
- ▶ $\forall k \in \mathbb{R} \setminus \{0\}$, we denote thus (kx, ky, k) as the homogeneous representation of the 2D point (x, y) .
- ▶ An arbitrary homogeneous $\mathbf{x} = (x_1, x_2, x_3)$ corresponds to the 2D point $(x_1/x_3, x_2/x_3)$.

Homogeneous representation of 2D lines and points

- ▶ A 2D line is defined by $ax + by + c = 0$ i.e. a parametrization $\mathbf{l} = (a, b, c)$.
- ▶ However, $kax + kby + kc = 0$ corresponds to the same line, thus $\mathbf{l} = (ka, kb, kc), \forall k \in \mathbb{R} \setminus \{0\}$
- ▶ A 2D point (x, y) lies on a line (a, b, c) if $ax + by + c = 0$.
- ▶ This may be expressed as $(x, y, 1)^T \cdot (a, b, c) = (x, y, 1)^T \cdot \mathbf{l} = 0$.
- ▶ $\forall k \in \mathbb{R} \setminus \{0\}, (kx, ky, k)^T \cdot \mathbf{l} = 0$ if and only if $(x, y, 1)^T \cdot \mathbf{l} = 0$.
- ▶ $\forall k \in \mathbb{R} \setminus \{0\}$, we denote thus (kx, ky, k) as the homogeneous representation of the 2D point (x, y) .
- ▶ An arbitrary homogeneous $\mathbf{x} = (x_1, x_2, x_3)$ corresponds to the 2D point $(x_1/x_3, x_2/x_3)$.
- ▶ **Result** : the point \mathbf{x} lies on the line \mathbf{l} if and only if $\mathbf{x}^T \mathbf{l} = 0$.

Homogeneous representation of 2D lines and points

- ▶ A 2D line is defined by $ax + by + c = 0$ i.e. a parametrization $\mathbf{l} = (a, b, c)$.
- ▶ However, $kax + kby + kc = 0$ corresponds to the same line, thus $\mathbf{l} = (ka, kb, kc), \forall k \in \mathbb{R} \setminus \{0\}$
- ▶ A 2D point (x, y) lies on a line (a, b, c) if $ax + by + c = 0$.
- ▶ This may be expressed as $(x, y, 1)^T \cdot (a, b, c) = (x, y, 1)^T \cdot \mathbf{l} = 0$.
- ▶ $\forall k \in \mathbb{R} \setminus \{0\}, (kx, ky, k)^T \cdot \mathbf{l} = 0$ if and only if $(x, y, 1)^T \cdot \mathbf{l} = 0$.
- ▶ $\forall k \in \mathbb{R} \setminus \{0\}$, we denote thus (kx, ky, k) as the homogeneous representation of the 2D point (x, y) .
- ▶ An arbitrary homogeneous $\mathbf{x} = (x_1, x_2, x_3)$ corresponds to the 2D point $(x_1/x_3, x_2/x_3)$.
- ▶ **Result** : the point \mathbf{x} lies on the line \mathbf{l} if and only if $\mathbf{x}^T \mathbf{l} = 0$.
- ▶ **Result** : the intersection of two lines \mathbf{l} and \mathbf{l}' is the point $\mathbf{x} = \mathbf{l} \times \mathbf{l}'$.

Homogeneous representation of 2D lines and points

- ▶ A 2D line is defined by $ax + by + c = 0$ i.e. a parametrization $\mathbf{l} = (a, b, c)$.
- ▶ However, $kax + kby + kc = 0$ corresponds to the same line, thus $\mathbf{l} = (ka, kb, kc), \forall k \in \mathbb{R} \setminus \{0\}$
- ▶ A 2D point (x, y) lies on a line (a, b, c) if $ax + by + c = 0$.
- ▶ This may be expressed as $(x, y, 1)^T \cdot (a, b, c) = (x, y, 1)^T \cdot \mathbf{l} = 0$.
- ▶ $\forall k \in \mathbb{R} \setminus \{0\}, (kx, ky, k)^T \cdot \mathbf{l} = 0$ if and only if $(x, y, 1)^T \cdot \mathbf{l} = 0$.
- ▶ $\forall k \in \mathbb{R} \setminus \{0\}$, we denote thus (kx, ky, k) as the homogeneous representation of the 2D point (x, y) .
- ▶ An arbitrary homogeneous $\mathbf{x} = (x_1, x_2, x_3)$ corresponds to the 2D point $(x_1/x_3, x_2/x_3)$.
- ▶ **Result** : the point \mathbf{x} lies on the line \mathbf{l} if and only if $\mathbf{x}^T \mathbf{l} = 0$.
- ▶ **Result** : the intersection of two lines \mathbf{l} and \mathbf{l}' is the point $\mathbf{x} = \mathbf{l} \times \mathbf{l}'$.
- ▶ **Result** : the line through two points \mathbf{x} and \mathbf{x}' is $\mathbf{l} = \mathbf{x} \times \mathbf{x}'$.

Some quick vector operations

$$\mathbf{x} \times \mathbf{y} = \mathbf{x}_{\times} \cdot \mathbf{y} = \begin{vmatrix} \mathbf{i} & \mathbf{j} & \mathbf{k} \\ x_1 & x_2 & x_3 \\ y_1 & y_2 & y_3 \end{vmatrix} = \begin{pmatrix} x_2 y_3 - x_3 y_2 \\ x_3 y_1 - x_1 y_3 \\ x_1 y_2 - y_1 x_2 \end{pmatrix}$$
$$\mathbf{x}_{\times} = \begin{pmatrix} 0 & -x_3 & x_2 \\ x_3 & 0 & -x_1 \\ -x_2 & x_1 & 0 \end{pmatrix}$$

Mixed product : $\mathbf{x}^T(\mathbf{y} \times \mathbf{z}) = |\mathbf{x} \ \mathbf{y} \ \mathbf{z}|$ (the volume of the parallelepiped defined by the three vectors)

Singular value decomposition

Theorem (SVD) :

Let \mathbf{A} be an $m \times n$ matrix. \mathbf{A} may be expressed as :

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \sum_{i=1}^{\min(m,n)} \sigma_i U_i V_i^T$$

where $\mathbf{\Sigma}$ is a $m \times n$ diagonal matrix with $\sigma_i = \mathbf{\Sigma}_{ii} \geq 0$, and \mathbf{U} ($m \times m$) and \mathbf{V} ($n \times n$) are composed of orthonormal columns

Singular value decomposition

Theorem (SVD) :

Let \mathbf{A} be an $m \times n$ matrix. \mathbf{A} may be expressed as :

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \sum_{i=1}^{\min(m,n)} \sigma_i U_i V_i^T$$

where $\mathbf{\Sigma}$ is a $m \times n$ diagonal matrix with $\sigma_i = \mathbf{\Sigma}_{ii} \geq 0$, and \mathbf{U} ($m \times m$) and \mathbf{V} ($n \times n$) are composed of orthonormal columns

- The rank of \mathbf{A} is the number of $\sigma_i > 0$

Singular value decomposition

Theorem (SVD) :

Let \mathbf{A} be an $m \times n$ matrix. \mathbf{A} may be expressed as :

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \sum_{i=1}^{\min(m,n)} \sigma_i U_i V_i^T$$

where $\mathbf{\Sigma}$ is a $m \times n$ diagonal matrix with $\sigma_i = \mathbf{\Sigma}_{ii} \geq 0$, and \mathbf{U} ($m \times m$) and \mathbf{V} ($n \times n$) are composed of orthonormal columns

- ▶ The rank of \mathbf{A} is the number of $\sigma_i > 0$
- ▶ An orthonormal basis for the null space of \mathbf{A} is composed of V_i for indices i such that $\sigma_i = 0$

Singular value decomposition

Theorem (SVD) :

Let \mathbf{A} be an $m \times n$ matrix. \mathbf{A} may be expressed as :

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \sum_{i=1}^{\min(m,n)} \sigma_i U_i V_i^T$$

where $\mathbf{\Sigma}$ is a $m \times n$ diagonal matrix with $\sigma_i = \mathbf{\Sigma}_{ii} \geq 0$, and \mathbf{U} ($m \times m$) and \mathbf{V} ($n \times n$) are composed of orthonormal columns

- ▶ The rank of \mathbf{A} is the number of $\sigma_i > 0$
- ▶ An orthonormal basis for the null space of \mathbf{A} is composed of V_i for indices i such that $\sigma_i = 0$
- ▶ By convention, the σ_i are aligned in descending order by the decomposition algorithms.

Outline

- The 3D representation of points
- The pinhole camera model
- Applying a coordinate transformation
- Homogeneous representations and algebraic operations
- The fundamental matrix
- The essential matrix
- Rectification

Why is this part “fundamental” ? (cheap joke)

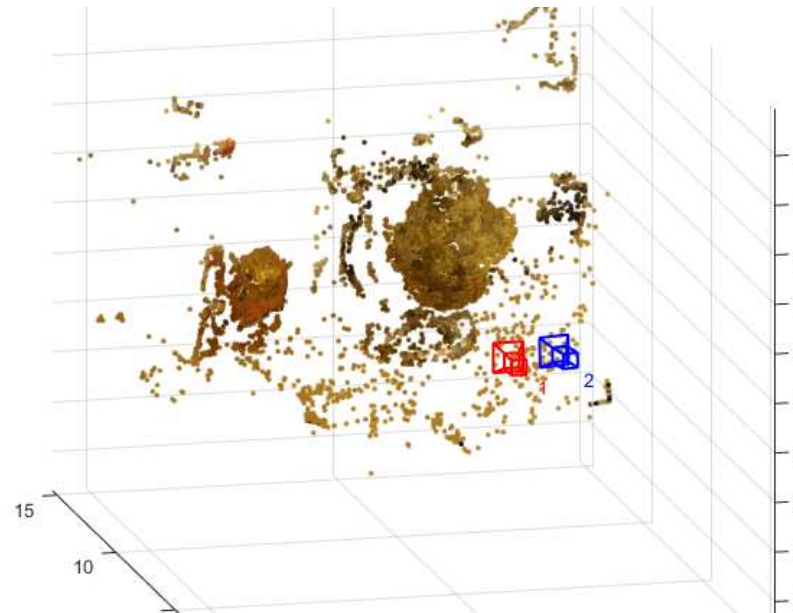
What we can get from two views :



Why is this part “fundamental” ? (cheap joke)

What we can get from two views :


- Sparse 3D reconstruction



Why is this part “fundamental” ? (cheap joke)

What we can get from two views :


- ▶ Sparse 3D reconstruction
- ▶ Relative camera pose estimation



Why is this part “fundamental” ? (cheap joke)

What we can get from two views :


- ▶ Sparse 3D reconstruction
- ▶ Relative camera pose estimation
- ▶ Parametric surface fitting



Why is this part “fundamental” ? (cheap joke)

What we can get from two views :


- ▶ Sparse 3D reconstruction
- ▶ Relative camera pose estimation
- ▶ Parametric surface fitting
- ▶ Dense 3D reconstruction (more complex work required for this)



Why is this part “fundamental” ? (cheap joke)

What we can get from two views :


- ▶ Sparse 3D reconstruction
- ▶ Relative camera pose estimation
- ▶ Parametric surface fitting
- ▶ Dense 3D reconstruction (more complex work required for this)
- ▶ ... but also many multi-view algorithms extend nicely from two-view analysis



The anatomy of two views

Some important observations :


- ▶ the pixel projection is along the ray defined by the 3D point and the camera center (i.e. as for \mathbf{x} , \mathbf{X} and \mathbf{C})



The anatomy of two views

Some important observations :


- ▶ the pixel projection is along the ray defined by the 3D point and the camera center (i.e. as for \mathbf{x} , \mathbf{X} and \mathbf{C})
- ▶ conversely, if \mathbf{x} and \mathbf{x}' do correspond to the same 3D point, the two rays intersect



The anatomy of two views

Some important observations :


- ▶ the pixel projection is along the ray defined by the 3D point and the camera center (i.e. as for \mathbf{x} , \mathbf{X} and \mathbf{C})
- ▶ conversely, if \mathbf{x} and \mathbf{x}' do correspond to the same 3D point, the two rays intersect
- ▶ the two rays define a plane π denoted as *epipolar plane*



The anatomy of two views

Some important observations :


- ▶ the pixel projection is along the ray defined by the 3D point and the camera center (i.e. as for \mathbf{x} , \mathbf{X} and \mathbf{C})
- ▶ conversely, if \mathbf{x} and \mathbf{x}' do correspond to the same 3D point, the two rays intersect
- ▶ the two rays define a plane π denoted as *epipolar plane*
- ▶ the epipolar plane also contains the ray defined by the camera centers



The anatomy of two views

From the projection in the two views we have :

$$\lambda \mathbf{x} = \mathbf{KX} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\mathbf{RX} + \mathbf{t})$$



The anatomy of two views

From the projection in the two views we have :

$$\lambda \mathbf{x} = \mathbf{KX} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\mathbf{RX} + \mathbf{t})$$

By eliminating \mathbf{X} we get :

$$\mathbf{x} = \lambda \mathbf{K}^{-1} \mathbf{x} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\lambda \mathbf{R} \mathbf{K}^{-1} \mathbf{x} + \mathbf{t})$$

The anatomy of two views

From the projection in the two views we have :

$$\lambda \mathbf{x} = \mathbf{KX} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\mathbf{RX} + \mathbf{t})$$

By eliminating \mathbf{X} we get :

$$\mathbf{x} = \lambda \mathbf{K}^{-1} \mathbf{x} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\lambda \mathbf{RK}^{-1} \mathbf{x} + \mathbf{t})$$

$$\lambda' \mathbf{K}'^{-1} \mathbf{x}' = \lambda \mathbf{RK}^{-1} \mathbf{x} + \mathbf{t}$$

The anatomy of two views

From the projection in the two views we have :

$$\lambda \mathbf{x} = \mathbf{KX} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\mathbf{RX} + \mathbf{t})$$

By eliminating \mathbf{X} we get :

$$\mathbf{X} = \lambda \mathbf{K}^{-1} \mathbf{x} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\lambda \mathbf{RK}^{-1} \mathbf{x} + \mathbf{t})$$

$$\lambda' \mathbf{K}'^{-1} \mathbf{x}' = \lambda \mathbf{RK}^{-1} \mathbf{x} + \mathbf{t}$$

We eliminate the sum by applying a cross product with \mathbf{t} :

$$\lambda' \mathbf{t} \times \mathbf{K}'^{-1} \mathbf{x}' = \lambda \mathbf{t} \times \mathbf{RK}^{-1} \mathbf{x}$$

The anatomy of two views

From the projection in the two views we have :

$$\lambda \mathbf{x} = \mathbf{KX} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\mathbf{RX} + \mathbf{t})$$

By eliminating \mathbf{X} we get :

$$\mathbf{x} = \lambda \mathbf{K}^{-1} \mathbf{x} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\lambda \mathbf{R} \mathbf{K}^{-1} \mathbf{x} + \mathbf{t})$$

$$\lambda' \mathbf{K}'^{-1} \mathbf{x}' = \lambda \mathbf{R} \mathbf{K}^{-1} \mathbf{x} + \mathbf{t}$$

We eliminate the sum by applying a cross product with \mathbf{t} :

$$\lambda' \mathbf{t} \times \mathbf{K}'^{-1} \mathbf{x}' = \lambda \mathbf{t} \times \mathbf{R} \mathbf{K}^{-1} \mathbf{x}$$

We multiply by $\mathbf{K}'^{-1} \mathbf{x}'$ in order to get a null mixed product :

$$0 = \lambda \mathbf{K}'^{-1} \mathbf{x}' \mathbf{t} \times \mathbf{R} \mathbf{K}^{-1} \mathbf{x}$$

The anatomy of two views

From the projection in the two views we have :

$$\lambda \mathbf{x} = \mathbf{KX} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\mathbf{RX} + \mathbf{t})$$

By eliminating \mathbf{X} we get :

$$\mathbf{x} = \lambda \mathbf{K}^{-1} \mathbf{x} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\lambda \mathbf{RK}^{-1} \mathbf{x} + \mathbf{t})$$

$$\lambda' \mathbf{K}'^{-1} \mathbf{x}' = \lambda \mathbf{RK}^{-1} \mathbf{x} + \mathbf{t}$$

We eliminate the sum by applying a cross product with \mathbf{t} :

$$\lambda' \mathbf{t} \times \mathbf{K}'^{-1} \mathbf{x}' = \lambda \mathbf{t} \times \mathbf{RK}^{-1} \mathbf{x}$$

We multiply by $\mathbf{K}'^{-1} \mathbf{x}'$ in order to get a null mixed product :

$$0 = \lambda \mathbf{K}'^{-1} \mathbf{x}' \mathbf{t} \times \mathbf{RK}^{-1} \mathbf{x}$$

Finally, by transposing $\mathbf{K}'^{-1} \mathbf{x}'$ and ignoring the scalar λ we get :

$$\mathbf{x}'^T \underbrace{\mathbf{K}'^{-T} \mathbf{t} \times \mathbf{RK}^{-1}}_{\mathbf{F}} \mathbf{x} = 0$$

The fundamental matrix \mathbf{F}

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$

- ▶ applying the \mathbf{F} constraint does not require information about the scene 3D structure

The fundamental matrix \mathbf{F}

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$

- ▶ applying the \mathbf{F} constraint does not require information about the scene 3D structure
- ▶ \mathbf{F} is valid for the whole image

The fundamental matrix \mathbf{F}

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$

- ▶ applying the \mathbf{F} constraint does not require information about the scene 3D structure
- ▶ \mathbf{F} is valid for the whole image
- ▶ we may apply the constraint without performing/knowning the camera calibration

The fundamental matrix \mathbf{F}

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$

- ▶ applying the \mathbf{F} constraint does not require information about the scene 3D structure
- ▶ \mathbf{F} is valid for the whole image
- ▶ we may apply the constraint without performing/knowning the camera calibration
- ▶ For a given point \mathbf{x}' , we denote by \mathbf{l}' its corresponding *epipolar line*. It follows from $\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$ that

$$\mathbf{l}' = \mathbf{F} \mathbf{x}$$

The fundamental matrix \mathbf{F}

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$

- ▶ applying the \mathbf{F} constraint does not require information about the scene 3D structure
- ▶ \mathbf{F} is valid for the whole image
- ▶ we may apply the constraint without performing/knowning the camera calibration
- ▶ For a given point \mathbf{x}' , we denote by \mathbf{l}' its corresponding *epipolar line*. It follows from $\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$ that

$$\mathbf{l}' = \mathbf{F} \mathbf{x}$$

- ▶ Similarly, $\mathbf{l} = \mathbf{F}^T \mathbf{x}'$

The fundamental matrix \mathbf{F}

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$

- ▶ applying the \mathbf{F} constraint does not require information about the scene 3D structure
- ▶ \mathbf{F} is valid for the whole image
- ▶ we may apply the constraint without performing/knowning the camera calibration
- ▶ For a given point \mathbf{x}' , we denote by \mathbf{l}' its corresponding *epipolar line*. It follows from $\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$ that

$$\mathbf{l}' = \mathbf{F} \mathbf{x}$$

- ▶ Similarly, $\mathbf{l} = \mathbf{F}^T \mathbf{x}'$
- ▶ The fundamental matrix constraint translates to a search along the epipolar line ...

The fundamental matrix \mathbf{F}

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$

- ▶ applying the \mathbf{F} constraint does not require information about the scene 3D structure
- ▶ \mathbf{F} is valid for the whole image
- ▶ we may apply the constraint without performing/knowning the camera calibration
- ▶ For a given point \mathbf{x}' , we denote by \mathbf{l}' its corresponding *epipolar line*. It follows from $\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$ that

$$\mathbf{l}' = \mathbf{F} \mathbf{x}$$

- ▶ Similarly, $\mathbf{l} = \mathbf{F}^T \mathbf{x}'$
- ▶ The fundamental matrix constraint translates to a search along the epipolar line ...
- ▶ ... but also $\mathbf{F} = \mathbf{K}'^{-T} \mathbf{t}_{\times} \mathbf{R} \mathbf{K}^{-1}$ encodes, along with the calibration matrices, *the rotation and translation* between views

The fundamental matrix \mathbf{F}

Theorem

The condition which is necessary and sufficient for a matrix \mathbf{F} to be a fundamental matrix is that

$$\det(\mathbf{F}) = 0$$

Multiple ways to notice that \mathbf{F} is rank deficient :

- ▶ it follows from the fact that $\det(\mathbf{t}_{\times}) = 0$

The fundamental matrix \mathbf{F}

Theorem

The condition which is necessary and sufficient for a matrix \mathbf{F} to be a fundamental matrix is that

$$\det(\mathbf{F}) = 0$$

Multiple ways to notice that \mathbf{F} is rank deficient :

- ▶ it follows from the fact that $\det(\mathbf{t}_{\times}) = 0$
- ▶ it follows from the fact that $\mathbf{F}\mathbf{e} = 0$

Computing \mathbf{F} - the 8 point algorithm

Straightforward approach :

- ▶ each observation (match) provides a constraint on \mathbf{F} as $\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$

Computing \mathbf{F} - the 8 point algorithm

Straightforward approach :

- ▶ each observation (match) provides a constraint on \mathbf{F} as $\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$
- ▶ if we group the unknowns as the column vector $\mathbf{f} = [f_{11} \ f_{12} \ \dots \ f_{33}]$, the constraint may be expressed as $\mathbf{a}_i \mathbf{f} = 0$, with \mathbf{a}_i a row vector

Computing F - the 8 point algorithm

Straightforward approach :

- ▶ each observation (match) provides a constraint on F as $\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$
- ▶ if we group the unknowns as the column vector $\mathbf{f} = [f_{11} \ f_{12} \ \dots \ f_{33}]$, the constraint may be expressed as $\mathbf{a}_i \mathbf{f} = 0$, with \mathbf{a}_i a row vector
- ▶ only 8 parameters are independent, since the scale is not determined

Computing \mathbf{F} - the 8 point algorithm

Straightforward approach :

- ▶ each observation (match) provides a constraint on \mathbf{F} as $\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$
- ▶ if we group the unknowns as the column vector $\mathbf{f} = [f_{11} \ f_{12} \ \dots \ f_{33}]$, the constraint may be expressed as $\mathbf{a}_i \mathbf{f} = 0$, with \mathbf{a}_i a row vector
- ▶ only 8 parameters are independent, since the scale is not determined
- ▶ the search for \mathbf{f} may be expressed as :

$$\min_{\mathbf{f}} \|\mathbf{A}\mathbf{f}\|, \text{ subject to } \|\mathbf{f}\| = 1$$

where $\mathbf{A} = [\mathbf{a}_1 \ \mathbf{a}_2 \ \dots \ \mathbf{a}_8]$

Computing \mathbf{F} - the 8 point algorithm

Straightforward approach :

- ▶ each observation (match) provides a constraint on \mathbf{F} as $\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$
- ▶ if we group the unknowns as the column vector $\mathbf{f} = [f_{11} \ f_{12} \ \dots \ f_{33}]$, the constraint may be expressed as $\mathbf{a}_i \mathbf{f} = 0$, with \mathbf{a}_i a row vector
- ▶ only 8 parameters are independent, since the scale is not determined
- ▶ the search for \mathbf{f} may be expressed as :

$$\min_{\mathbf{f}} \|\mathbf{A}\mathbf{f}\|, \text{ subject to } \|\mathbf{f}\| = 1$$

where $\mathbf{A} = [\mathbf{a}_1 \ \mathbf{a}_2 \ \dots \ \mathbf{a}_8]$

- ▶ **Solution** : \mathbf{f} is the last column of \mathbf{V} , where $\mathbf{A} = \mathbf{U}\mathbf{D}\mathbf{V}^T$ is the SVD of \mathbf{A}

Computing \mathbf{F} - the 8 point algorithm

Straightforward approach :

- ▶ each observation (match) provides a constraint on \mathbf{F} as $\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$
- ▶ if we group the unknowns as the column vector $\mathbf{f} = [f_{11} \ f_{12} \ \dots \ f_{33}]$, the constraint may be expressed as $\mathbf{a}_i \mathbf{f} = 0$, with \mathbf{a}_i a row vector
- ▶ only 8 parameters are independent, since the scale is not determined
- ▶ the search for \mathbf{f} may be expressed as :

$$\min_{\mathbf{f}} \|\mathbf{A}\mathbf{f}\|, \text{ subject to } \|\mathbf{f}\| = 1$$

where $\mathbf{A} = [\mathbf{a}_1 \ \mathbf{a}_2 \ \dots \ \mathbf{a}_8]$

- ▶ **Solution** : \mathbf{f} is the last column of \mathbf{V} , where $\mathbf{A} = \mathbf{U}\mathbf{D}\mathbf{V}^T$ is the SVD of \mathbf{A}

- ▶ **Proof** :

$\|\mathbf{U}\mathbf{D}\mathbf{V}^T \mathbf{f}\| = \|\mathbf{D}\mathbf{V}^T \mathbf{f}\|$, and $\|\mathbf{f}\| = \|\mathbf{V}^T \mathbf{f}\|$. We have to minimize $\|\mathbf{D}\mathbf{V}^T \mathbf{f}\|$ subject to $\|\mathbf{V}^T \mathbf{f}\| = 1$. If $\mathbf{y} = \mathbf{V}^T \mathbf{f}$, then we minimize $\|\mathbf{D}\mathbf{y}\|$ subject to $\|\mathbf{y}\| = 1$. Since \mathbf{D} is diagonal with values in descending order, it means that $\mathbf{y} = (0, 0, \dots, 1)$, and $\mathbf{f} = \mathbf{V}\mathbf{y}$ is the last column of \mathbf{V} . (A5.3, Hartley and Zisserman)

Considerations - the 8 point algorithm

Straightforward approach :

- ▶ major issue : the solution **F** may violate the rank constraint !

Considerations - the 8 point algorithm

Straightforward approach :

- ▶ major issue : the solution **F** may violate the rank constraint !
- ▶ Hack : decompose **F** using SVD, set $\sigma_3 = 0$ and recompose.

Considerations - the 8 point algorithm

Straightforward approach :

- ▶ major issue : the solution **F** may violate the rank constraint !
- ▶ Hack : decompose **F** using SVD, set $\sigma_3 = 0$ and recompose.
- ▶ What about searching directly for a rank 2 solution for **F** ?

Considerations - the 8 point algorithm

Straightforward approach :

- ▶ major issue : the solution \mathbf{F} may violate the rank constraint !
- ▶ Hack : decompose \mathbf{F} using SVD, set $\sigma_3 = 0$ and recompose.
- ▶ What about searching directly for a rank 2 solution for \mathbf{F} ?

The 7 point algorithm :

- ▶ Use 7 constraints for $\mathbf{A}\mathbf{f} = \mathbf{0}$

Considerations - the 8 point algorithm

Straightforward approach :

- ▶ major issue : the solution \mathbf{F} may violate the rank constraint !
- ▶ Hack : decompose \mathbf{F} using SVD, set $\sigma_3 = 0$ and recompose.
- ▶ What about searching directly for a rank 2 solution for \mathbf{F} ?

The 7 point algorithm :

- ▶ Use 7 constraints for $\mathbf{A}\mathbf{f} = \mathbf{0}$
- ▶ Use SVD on \mathbf{A} in order to find the vectors \mathbf{f}_1 and \mathbf{f}_2 that span the null space (the kernel) of \mathbf{A}

Considerations - the 8 point algorithm

Straightforward approach :

- ▶ major issue : the solution \mathbf{F} may violate the rank constraint !
- ▶ Hack : decompose \mathbf{F} using SVD, set $\sigma_3 = 0$ and recompose.
- ▶ What about searching directly for a rank 2 solution for \mathbf{F} ?

The 7 point algorithm :

- ▶ Use 7 constraints for $\mathbf{A}\mathbf{f} = \mathbf{0}$
- ▶ Use SVD on \mathbf{A} in order to find the vectors \mathbf{f}_1 and \mathbf{f}_2 that span the null space (the kernel) of \mathbf{A}
- ▶ Find an element in the kernel expressed by the linear combination $\mathbf{f} = \mathbf{f}_1 + \alpha\mathbf{f}_2$ which also satisfies $\det(\mathbf{F}) = 0$

Considerations - the 8 point algorithm

Straightforward approach :

- ▶ major issue : the solution \mathbf{F} may violate the rank constraint !
- ▶ Hack : decompose \mathbf{F} using SVD, set $\sigma_3 = 0$ and recompose.
- ▶ What about searching directly for a rank 2 solution for \mathbf{F} ?

The 7 point algorithm :

- ▶ Use 7 constraints for $\mathbf{A}\mathbf{f} = \mathbf{0}$
- ▶ Use SVD on \mathbf{A} in order to find the vectors \mathbf{f}_1 and \mathbf{f}_2 that span the null space (the kernel) of \mathbf{A}
- ▶ Find an element in the kernel expressed by the linear combination $\mathbf{f} = \mathbf{f}_1 + \alpha\mathbf{f}_2$ which also satisfies $\det(\mathbf{F}) = 0$
- ▶ $\det(\mathbf{F}_1 + \alpha\mathbf{F}_2)$ is a third degree polynomial, so up to three potential solutions may be recovered

Considerations - the 8 point algorithm

Straightforward approach :

- ▶ major issue : the solution \mathbf{F} may violate the rank constraint !
- ▶ Hack : decompose \mathbf{F} using SVD, set $\sigma_3 = 0$ and recompose.
- ▶ What about searching directly for a rank 2 solution for \mathbf{F} ?

The 7 point algorithm :

- ▶ Use 7 constraints for $\mathbf{A}\mathbf{f} = \mathbf{0}$
- ▶ Use SVD on \mathbf{A} in order to find the vectors \mathbf{f}_1 and \mathbf{f}_2 that span the null space (the kernel) of \mathbf{A}
- ▶ Find an element in the kernel expressed by the linear combination $\mathbf{f} = \mathbf{f}_1 + \alpha\mathbf{f}_2$ which also satisfies $\det(\mathbf{F}) = 0$
- ▶ $\det(\mathbf{F}_1 + \alpha\mathbf{F}_2)$ is a third degree polynomial, so up to three potential solutions may be recovered
- ▶ This algorithm is also preferred as fewer observations are needed

Outline

- The 3D representation of points
- The pinhole camera model
- Applying a coordinate transformation
- Homogeneous representations and algebraic operations
- The fundamental matrix
- The essential matrix
- Rectification

Using the camera calibration and the essential matrix

If the calibration matrices \mathbf{K} and \mathbf{K}' are known :

- ▶ we may recover the pose information from $\mathbf{F} = \mathbf{K}'^{-T} \mathbf{t}_{\times} \mathbf{R} \mathbf{K}^{-1}$:

$$\mathbf{E} = \mathbf{t}_{\times} \mathbf{R} = \mathbf{K}'^T \mathbf{F} \mathbf{K}$$

Using the camera calibration and the essential matrix

If the calibration matrices \mathbf{K} and \mathbf{K}' are known :

- ▶ we may recover the pose information from $\mathbf{F} = \mathbf{K}'^{-T} \mathbf{t}_{\times} \mathbf{R} \mathbf{K}^{-1}$:

$$\mathbf{E} = \mathbf{t}_{\times} \mathbf{R} = \mathbf{K}'^T \mathbf{F} \mathbf{K}$$

- ▶ \mathbf{E} has five degrees of freedom (and not six) because the relative translation \mathbf{t} has a scale ambiguity (just as \mathbf{F}).

Using the camera calibration and the essential matrix

If the calibration matrices \mathbf{K} and \mathbf{K}' are known :

- ▶ we may recover the pose information from $\mathbf{F} = \mathbf{K}'^{-T} \mathbf{t}_{\times} \mathbf{R} \mathbf{K}^{-1}$:

$$\mathbf{E} = \mathbf{t}_{\times} \mathbf{R} = \mathbf{K}'^T \mathbf{F} \mathbf{K}$$

- ▶ \mathbf{E} has five degrees of freedom (and not six) because the relative translation \mathbf{t} has a scale ambiguity (just as \mathbf{F}).
- ▶ Beside $\det(\mathbf{E}) = 0$, there is an additional constraint with respect to \mathbf{F} , which results from the structure of \mathbf{E} :

Theorem : The condition which is necessary and sufficient for a matrix \mathbf{E} to be an essential matrix is that two of its singular values be equal, and the third one be 0.

Using the camera calibration and the essential matrix

If the calibration matrices \mathbf{K} and \mathbf{K}' are known :

- ▶ we may recover the pose information from $\mathbf{F} = \mathbf{K}'^{-T} \mathbf{t}_{\times} \mathbf{R} \mathbf{K}^{-1}$:

$$\mathbf{E} = \mathbf{t}_{\times} \mathbf{R} = \mathbf{K}'^T \mathbf{F} \mathbf{K}$$

- ▶ \mathbf{E} has five degrees of freedom (and not six) because the relative translation \mathbf{t} has a scale ambiguity (just as \mathbf{F}).
- ▶ Beside $\det(\mathbf{E}) = 0$, there is an additional constraint with respect to \mathbf{F} , which results from the structure of \mathbf{E} :

Theorem : The condition which is necessary and sufficient for a matrix \mathbf{E} to be an essential matrix is that two of its singular values be equal, and the third one be 0.

- ▶ There are thus at least five points needed for recovering directly \mathbf{E} from an image pair, assuming that the calibration matrices are known, and there is an algorithm which solves this minimal problem(Nistér, David. "An efficient solution to the five-point relative pose problem." IEEE Transactions on Pattern Analysis and Machine Intelligence (2004).)

Using the camera calibration and the essential matrix

If the calibration matrices \mathbf{K} and \mathbf{K}' are known :

- ▶ we may recover the pose information from $\mathbf{F} = \mathbf{K}'^{-T} \mathbf{t}_{\times} \mathbf{R} \mathbf{K}^{-1}$:

$$\mathbf{E} = \mathbf{t}_{\times} \mathbf{R} = \mathbf{K}'^T \mathbf{F} \mathbf{K}$$

- ▶ \mathbf{E} has five degrees of freedom (and not six) because the relative translation \mathbf{t} has a scale ambiguity (just as \mathbf{F}).
- ▶ Beside $\det(\mathbf{E}) = 0$, there is an additional constraint with respect to \mathbf{F} , which results from the structure of \mathbf{E} :

Theorem : The condition which is necessary and sufficient for a matrix \mathbf{E} to be an essential matrix is that two of its singular values be equal, and the third one be 0.

- ▶ There are thus at least five points needed for recovering directly \mathbf{E} from an image pair, assuming that the calibration matrices are known, and there is an algorithm which solves this minimal problem(Nistér, David. "An efficient solution to the five-point relative pose problem." IEEE Transactions on Pattern Analysis and Machine Intelligence (2004).)
- ▶ Knowing \mathbf{E} : interesting for relative pose estimation

Using the camera calibration and the essential matrix

If the calibration matrices \mathbf{K} and \mathbf{K}' are known :

- ▶ we may recover the pose information from $\mathbf{F} = \mathbf{K}'^{-T} \mathbf{t}_{\times} \mathbf{R} \mathbf{K}^{-1}$:

$$\mathbf{E} = \mathbf{t}_{\times} \mathbf{R} = \mathbf{K}'^T \mathbf{F} \mathbf{K}$$

- ▶ \mathbf{E} has five degrees of freedom (and not six) because the relative translation \mathbf{t} has a scale ambiguity (just as \mathbf{F}).
- ▶ Beside $\det(\mathbf{E}) = 0$, there is an additional constraint with respect to \mathbf{F} , which results from the structure of \mathbf{E} :

Theorem : The condition which is necessary and sufficient for a matrix \mathbf{E} to be an essential matrix is that two of its singular values be equal, and the third one be 0.

- ▶ There are thus at least five points needed for recovering directly \mathbf{E} from an image pair, assuming that the calibration matrices are known, and there is an algorithm which solves this minimal problem(Nistér, David. "An efficient solution to the five-point relative pose problem." IEEE Transactions on Pattern Analysis and Machine Intelligence (2004).)
- ▶ Knowing \mathbf{E} : interesting for relative pose estimation
- ▶ Main disadvantage : \mathbf{K} and \mathbf{K}' are required to get to \mathbf{E}

Recovering R and t from E

It has been shown that the decomposition of E is possible and there are actually four valid solutions (9.6.2, *Hartley and Zisserman*) :




Fig. 9.12. The four possible solutions for calibrated reconstruction from E . Between the left and right sides there is a baseline reversal. Between the top and bottom rows camera B rotates 180° about the baseline. Note, only in (a) is the reconstructed point in front of both cameras.

- Identify the correct solution : cheirality check (the 3D points have to be in front of the camera) with an additional match from the two views

Outline


- The 3D representation of points
- The pinhole camera model
- Applying a coordinate transformation
- Homogeneous representations and algebraic operations
- The fundamental matrix
- The essential matrix
- Rectification

Rectification

Using \mathbf{F} , we restrict the search for the corresponding projection \mathbf{x}' of a point \mathbf{x} to a line (the epipolar line $\mathbf{l}' = \mathbf{F}\mathbf{x}$).

Stereo rectification

- Apply an adjustment to the images in order to get horizontal epipolar lines in both views




Rectification

Using \mathbf{F} , we restrict the search for the corresponding projection \mathbf{x}' of a point \mathbf{x} to a line (the epipolar line $\mathbf{l}' = \mathbf{F}\mathbf{x}$).

Stereo rectification

- ▶ Apply an adjustment to the images in order to get horizontal epipolar lines in both views
- ▶ The search for \mathbf{x}' takes place simply along the same corresponding row in the second image : interesting for dense correspondence




Rectification

Using \mathbf{F} , we restrict the search for the corresponding projection \mathbf{x}' of a point \mathbf{x} to a line (the epipolar line $\mathbf{l}' = \mathbf{F}\mathbf{x}$).

Stereo rectification

- ▶ Apply an adjustment to the images in order to get horizontal epipolar lines in both views
- ▶ The search for \mathbf{x}' takes place simply along the same corresponding row in the second image : interesting for dense correspondence
- ▶ This implies that epipoles are at horizontal infinity : $\mathbf{e} = \mathbf{e}' = [1 \ 0 \ 0]^T$




Rectification

Using \mathbf{F} , we restrict the search for the corresponding projection \mathbf{x}' of a point \mathbf{x} to a line (the epipolar line $\mathbf{l}' = \mathbf{F}\mathbf{x}$).

Stereo rectification

- ▶ Apply an adjustment to the images in order to get horizontal epipolar lines in both views
- ▶ The search for \mathbf{x}' takes place simply along the same corresponding row in the second image : interesting for dense correspondence
- ▶ This implies that epipoles are at horizontal infinity : $\mathbf{e} = \mathbf{e}' = [1 \ 0 \ 0]^T$
- ▶ Apply a virtual rotation of cameras (Fusiello, A. ; Trucco, E. ; Verri, A. A compact algorithm for rectification of stereo pairs. Mach. Vision Appl 2000)




Rectification

Using \mathbf{F} , we restrict the search for the corresponding projection \mathbf{x}' of a point \mathbf{x} to a line (the epipolar line $\mathbf{l}' = \mathbf{F}\mathbf{x}$).


Stereo rectification

- ▶ Apply an adjustment to the images in order to get horizontal epipolar lines in both views
- ▶ The search for \mathbf{x}' takes place simply along the same corresponding row in the second image : interesting for dense correspondence
- ▶ This implies that epipoles are at horizontal infinity : $\mathbf{e} = \mathbf{e}' = [1\ 0\ 0]^T$
- ▶ Apply a virtual rotation of cameras (Fusiello, A. ; Trucco, E. ; Verri, A. A compact algorithm for rectification of stereo pairs. Mach. Vision Appl 2000)
- ▶ An interpolation is required for creating the new images, but high computation gain overall



Triangulation - the building block of 3D reprojections

We have the pose \mathbf{R}, \mathbf{t}' between cameras and the projection locations \mathbf{x}, \mathbf{x}' . What now ?



Get \mathbf{X} : triangulate the point in 3D

Triangulation - the building block of 3D reprojections

We have the pose \mathbf{R}, \mathbf{t}' between cameras and the projection locations \mathbf{x}, \mathbf{x}' . What now ?

Get \mathbf{X} : triangulate the point in 3D

► Back to our stereo projection equations :

$$\lambda \mathbf{x} = \mathbf{KX} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\mathbf{RX} + \mathbf{t})$$

Triangulation - the building block of 3D reprojections

We have the pose \mathbf{R}, \mathbf{t}' between cameras and the projection locations \mathbf{x}, \mathbf{x}' . What now ?

Get \mathbf{X} : triangulate the point in 3D

- ▶ Back to our stereo projection equations :

$$\lambda \mathbf{x} = \mathbf{KX} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\mathbf{RX} + \mathbf{t})$$

- ▶ We have five scalar unknowns and six equations - a direct approach is possible by solving an overdetermined linear system

Triangulation - the building block of 3D reprojections

We have the pose \mathbf{R}, \mathbf{t}' between cameras and the projection locations \mathbf{x}, \mathbf{x}' . What now ?

Get \mathbf{X} : triangulate the point in 3D

- ▶ Back to our stereo projection equations :

$$\lambda \mathbf{x} = \mathbf{KX} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\mathbf{RX} + \mathbf{t})$$

- ▶ We have five scalar unknowns and six equations - a direct approach is possible by solving an overdetermined linear system
- ▶ There are other algorithms which are more accurate, but costlier
Hartley, R. I., Sturm, P. (1997). Triangulation. Computer vision and image understanding, 68(2), 146-157
Lindstrom, Peter. "Triangulation made easy." In Computer Vision and Pattern Recognition (CVPR), 2010 IEEE Conference on, pp. 1554-1561

Triangulation - the building block of 3D reprojections

We have the pose \mathbf{R}, \mathbf{t}' between cameras and the projection locations \mathbf{x}, \mathbf{x}' . What now ?

Get \mathbf{X} : triangulate the point in 3D

- ▶ Back to our stereo projection equations :


$$\lambda \mathbf{x} = \mathbf{KX} \quad \lambda' \mathbf{x}' = \mathbf{K}'(\mathbf{RX} + \mathbf{t})$$

- ▶ We have five scalar unknowns and six equations - a direct approach is possible by solving an overdetermined linear system
- ▶ There are other algorithms which are more accurate, but costlier
Hartley, R. I., Sturm, P. (1997). Triangulation. Computer vision and image understanding, 68(2), 146-157
Lindstrom, Peter. "Triangulation made easy." In Computer Vision and Pattern Recognition (CVPR), 2010 IEEE Conference on, pp. 1554-1561
- ▶ The linear approach is reasonably good, and it is effective especially if used as an initialization for a nonlinear refinement (as we will see in the following slides)

Triangulation - how to use multiple views

If we have multiple views, the unknown \mathbf{X}_j may be constrained by multiple observations $\mathbf{z}_{j,\tau}$ from cameras C_τ characterized by some pose parametrization \mathbf{s}_τ . How to use them effectively together?

Nonlinear optimization




Triangulation - how to use multiple views

If we have multiple views, the unknown \mathbf{X}_j may be constrained by multiple observations $\mathbf{z}_{j,\tau}$ from cameras C_τ characterized by some pose parametrization \mathbf{s}_τ . How to use them effectively together?

Nonlinear optimization

- Analytical solutions are not practical, in most cases we solve the optimization iteratively



Triangulation - how to use multiple views


If we have multiple views, the unknown \mathbf{X}_j may be constrained by multiple observations $\mathbf{z}_{j,\tau}$ from cameras C_τ characterized by some pose parametrization \mathbf{s}_τ . How to use them effectively together?

Nonlinear optimization

- ▶ Analytical solutions are not practical, in most cases we solve the optimization iteratively
- ▶ We define an error related to each of the observation, i.e. the distance between the observation and the projection of \mathbf{X}_j : $e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_j) = \mathbf{z}_j - g(\mathbf{s}_\tau, \mathbf{X}_j)$, where g is the camera projection function. Then, we have :

$$\hat{\mathbf{X}}_j = \arg \min_{\mathbf{X}_j} \sum_{\tau} e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_j)^T e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_j)$$

- ▶ Use Gauss-Newton or LM (usually the optimum is not far from a reasonable initialization)



Triangulation - how to use multiple views


If we have multiple views, the unknown \mathbf{X}_j may be constrained by multiple observations $\mathbf{z}_{j,\tau}$ from cameras C_τ characterized by some pose parametrization \mathbf{s}_τ . How to use them effectively together?

Nonlinear optimization

- ▶ Analytical solutions are not practical, in most cases we solve the optimization iteratively
- ▶ We define an error related to each of the observation, i.e. the distance between the observation and the projection of \mathbf{X}_j : $e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_j) = \mathbf{z}_j - g(\mathbf{s}_\tau, \mathbf{X}_j)$, where g is the camera projection function. Then, we have :

$$\hat{\mathbf{X}}_j = \arg \min_{\mathbf{X}_j} \sum_{\tau} e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_j)^T e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_j)$$


- ▶ Use Gauss-Newton or LM (usually the optimum is not far from a reasonable initialization)
- ▶ More than one 3D point may be refined, but in this way the optimizations are decoupled



Pose estimation - how to use multiple views

Opposite problem : we have a set of 3D points \mathbf{X}_j (computed previously) which are visible from camera C_τ . Based on current observations $\mathbf{z}_{j,\tau}$ from C_τ we would like to estimate its pose \mathbf{s}_τ .

Nonlinear optimization




Pose estimation - how to use multiple views

Opposite problem : we have a set of 3D points \mathbf{X}_j (computed previously) which are visible from camera C_τ . Based on current observations $\mathbf{z}_{j,\tau}$ from C_τ we would like to estimate its pose \mathbf{s}_τ .

Nonlinear optimization

- We define an error related to each of the observations, i.e. the distance between the observation and the projection of \mathbf{X}_j : $e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau}) = \mathbf{z}_{j,\tau} - g(\mathbf{s}_\tau, \mathbf{X}_j)$, where g is the camera projection function. Then, we have :

$$\hat{\mathbf{s}}_\tau = \arg \min_{\mathbf{s}_\tau} \sum_j e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau})^T e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau})$$



Pose estimation - how to use multiple views


Opposite problem : we have a set of 3D points \mathbf{X}_j (computed previously) which are visible from camera C_τ . Based on current observations $\mathbf{z}_{j,\tau}$ from C_τ we would like to estimate its pose \mathbf{s}_τ .

Nonlinear optimization

- ▶ We define an error related to each of the observations, i.e. the distance between the observation and the projection of \mathbf{X}_j : $e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau}) = \mathbf{z}_{j,\tau} - g(\mathbf{s}_\tau, \mathbf{X}_j)$, where g is the camera projection function. Then, we have :

$$\hat{\mathbf{s}}_\tau = \arg \min_{\mathbf{s}_\tau} \sum_j e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau})^T e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau})$$

- ▶ Use Gauss-Newton or LM, but the initialization is very important. Two strategies help :



Pose estimation - how to use multiple views


Opposite problem : we have a set of 3D points \mathbf{X}_j (computed previously) which are visible from camera C_τ . Based on current observations $\mathbf{z}_{j,\tau}$ from C_τ we would like to estimate its pose \mathbf{s}_τ .

Nonlinear optimization

- ▶ We define an error related to each of the observations, i.e. the distance between the observation and the projection of \mathbf{X}_j : $e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau}) = \mathbf{z}_{j,\tau} - g(\mathbf{s}_\tau, \mathbf{X}_j)$, where g is the camera projection function. Then, we have :

$$\hat{\mathbf{s}}_\tau = \arg \min_{\mathbf{s}_\tau} \sum_j e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau})^T e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau})$$

- ▶ Use Gauss-Newton or LM, but the initialization is very important. Two strategies help :
 - ▶ if the camera is moving, predict the current location based on its previous trajectory



Pose estimation - how to use multiple views


Opposite problem : we have a set of 3D points \mathbf{X}_j (computed previously) which are visible from camera C_τ . Based on current observations $\mathbf{z}_{j,\tau}$ from C_τ we would like to estimate its pose \mathbf{s}_τ .

Nonlinear optimization

- ▶ We define an error related to each of the observations, i.e. the distance between the observation and the projection of \mathbf{X}_j : $e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau}) = \mathbf{z}_{j,\tau} - g(\mathbf{s}_\tau, \mathbf{X}_j)$, where g is the camera projection function. Then, we have :

$$\hat{\mathbf{s}}_\tau = \arg \min_{\mathbf{s}_\tau} \sum_j e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau})^T e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau})$$

- ▶ Use Gauss-Newton or LM, but the initialization is very important. Two strategies help :
 - ▶ if the camera is moving, predict the current location based on its previous trajectory
 - ▶ from the projection of three 3D points in space and their projections, one may compute the camera pose in a closed form (the P3P problem)



Limitations of previous approaches

Assumptions :

Limitations of previous approaches

Assumptions :

- ▶ for triangulation : we assume that the pose is correctly estimated

Limitations of previous approaches

Assumptions :

- ▶ for triangulation : we assume that the pose is correctly estimated
- ▶ for pose estimation : we assume that the 3D locations are accurate

Limitations of previous approaches

Assumptions :

- ▶ for triangulation : we assume that the pose is correctly estimated
- ▶ for pose estimation : we assume that the 3D locations are accurate
- ▶ in reality all estimations we perform are noisy

Limitations of previous approaches

Assumptions :

- ▶ for triangulation : we assume that the pose is correctly estimated
- ▶ for pose estimation : we assume that the 3D locations are accurate
- ▶ in reality all estimations we perform are noisy
- ▶ if we also apply the process iteratively (triangulation, pose estimation and repeat) the errors will be amplified (drift)

Global optimization - initial step

Since computational power is widely available for autonomous systems, we favour a solution which minimizes jointly with respect to the point locations and to the poses.

Initial step :

Global optimization - initial step

Since computational power is widely available for autonomous systems, we favour a solution which minimizes jointly with respect to the point locations and to the poses.

Initial step :

- ▶ we will just add a new unknown pose to the previous set of variables and refine it :

$$\hat{\mathbf{s}}_{\tau} = \arg \min_{\mathbf{s}_{\tau}} \sum_j e(\mathbf{s}_{\tau}, \mathbf{X}_j, \mathbf{z}_{j,\tau})^T e(\mathbf{s}_{\tau}, \mathbf{X}_j, \mathbf{z}_{j,\tau})$$

Global optimization - initial step

Since computational power is widely available for autonomous systems, we favour a solution which minimizes jointly with respect to the point locations and to the poses.

Initial step :

- ▶ we will just add a new unknown pose to the previous set of variables and refine it :

$$\hat{\mathbf{s}}_{\tau} = \arg \min_{\mathbf{s}_{\tau}} \sum_j e(\mathbf{s}_{\tau}, \mathbf{X}_j, \mathbf{z}_{j,\tau})^T e(\mathbf{s}_{\tau}, \mathbf{X}_j, \mathbf{z}_{j,\tau})$$

- ▶ observation : this step does not modify \mathbf{X}

Global optimization - initial step

Since computational power is widely available for autonomous systems, we favour a solution which minimizes jointly with respect to the point locations and to the poses.

Initial step :

- ▶ we will just add a new unknown pose to the previous set of variables and refine it :

$$\hat{\mathbf{s}}_\tau = \arg \min_{\mathbf{s}_\tau} \sum_j e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau})^T e(\mathbf{s}_\tau, \mathbf{X}_j, \mathbf{z}_{j,\tau})$$

- ▶ observation : this step does not modify \mathbf{X}
- ▶ the interest of the initial step is just to provide a quality initialization for \mathbf{s}_τ as $\hat{\mathbf{s}}_t$


Global optimization - final step

We compute the MAP (Maximum A Posteriori) for the maximum amount of preliminary estimations and observations that we have at that moment (brutal, massive optimization). The solution we search this time is provided by :

$$\tilde{\mathbf{S}}_{0:t}, \tilde{\mathbf{X}} = \arg \min_{\mathbf{S}_{0:t}, \mathbf{X}} \sum_{\tau=0}^T \sum_{j=1}^M e(\mathbf{s}_{\tau}, \mathbf{X}_{j,\tau}, z_{j,\tau})^T e(\mathbf{s}_{\tau}, \mathbf{X}_{j,\tau}, z_{j,\tau})$$

The complexity of this algorithm, once we exploit the sparseness of its Jacobian : $O(T^3 + MT^2)$, which is very interesting since $M \gg T$.

Towards real time reconstruction



An example of configuration : 5207 3D points, 54 poses, 24609 projections, 15945 variables, 21 it., 7.99 sec.

Not fast enough !

- ▶ Selection of key-frames
- ▶ Parallel execution of tracking et BA (initial and final steps)
- ▶ Limit the number of iterations (when needed)
- ▶ Local Bundle Adjustment

Typical architecture for RT optimization

